

Deep Learning Project2 Report

Chenkai Wang
323452
chenkai.wang@epfl.ch

Siran Li
321825
siran.li@epfl.ch

Shijian Xu
327448
shijian.xu@epfl.ch

Abstract—The objective of this project is to implement a mini “deep learning framework” using only PyTorch’s tensor operations and the standard math library. The task is to use this mini framework to train a Multi-Layer Perceptron (MLP) for a binary classification problem. To achieve this goal, we implement the modules Linear, ReLU, Tanh, Sigmoid, SELU, Sequential, MSELoss and SGD optimizer. The interfaces of these modules are very close to their PyTorch implementation.

I. IMPLEMENTATION

The mini deep learning framework consists of several modules: Linear, Sequential, MSELoss, SGD optimizer and several activation functions, which are ReLU, Tanh, Sigmoid and SELU. All of these modules, except the SGD optimizer, inherit from a basic Module class. The final MLP model also inherits from the Module class and is composed of a Sequential member.

A. Data Description

In order to test our framework and to see its performance, we generate a training and a test dataset of 1000 points separately, uniformly sampled from $[0, 1]^2$. Each point has a label 0 if it is outside the disk centered at $(0.5, 0.5)$ of radius $1/\sqrt{2\pi}$, and 1 inside. The visualization of the data point distribution is shown in Fig. 1.




Fig. 1: Visualization of data set.

B. Modules

1) *Module*: The base class Module defines the functions each module needs to implement. The code snippet of the implementation of the basic Module class is shown below.

```
1 class Module(object):
2     """
3     Base Class. Other modules should inherit from it
4     and rewrite the 'forward', 'backward' and
5     'param' functions.
```

```
"""
6     def __call__(self, *input):
7         return self.forward(*input)
8
9     def forward(self, *input):
10        raise NotImplementedError
11
12    def backward(self, *gradwrtoutput):
13        raise NotImplementedError
14
15    def parameters(self):
16        return []
```

Listing 1: The Base Module Class

We also use the Python built-in function `__call__` to make the modules’ interfaces behave more PyTorch-like. In this case, we can call the module directly (`model(x)`) instead of `model.forward(x)`.

2) *Linear*: Linear layer is the core in the framework. In our implementation, this layer takes in a 2-d tensor $x \in \mathbb{R}^{N \times in_features}$, where N is the batch size. The forward output is produced by

$$y = xW + b \quad (1)$$

where W is the weight matrix with shape $in_features \times out_features$, and b is the bias vector. Both W and b are initialized with uniform distribution from $\mathcal{U}(-\sqrt{k}, \sqrt{k})$, where $k = \frac{1}{in_features}$. The backward function takes in the gradient of loss with respect to the output y , $\frac{\partial \ell}{\partial y}$, and compute the gradient for W, b and x :

$$\begin{aligned} \frac{\partial \ell}{\partial W} &= x^T \frac{\partial \ell}{\partial y} \\ \frac{\partial \ell}{\partial b} &= \frac{\partial \ell}{\partial y} \\ \frac{\partial \ell}{\partial x} &= \frac{\partial \ell}{\partial y} W^T \end{aligned} \quad (2)$$

The gradient of x will be returned and propagated to the former layer.

3) Activation Functions:

- **ReLU** The ReLU activation takes in the feature x and applies the following function element-wise:

$$y = ReLU(x) = (x)^+ = \begin{cases} x & \text{if } x > 0 \\ 0 & \text{if } x \leq 0 \end{cases} \quad (3)$$

The gradient of ReLU is:

$$\frac{\partial y}{\partial x} = \begin{cases} 1 & \text{if } x > 0 \\ 0 & \text{if } x \leq 0 \end{cases} \quad (4)$$

- **Tanh** The Tanh activation takes in the feature x and applies the following function element-wise:

$$y = \text{Tanh}(x) = \frac{\exp(x) - \exp(-x)}{\exp(x) + \exp(-x)} \quad (5)$$

The gradient of Tanh is :

$$\frac{\partial y}{\partial x} = 1 - \text{Tanh}(x)^2 \quad (6)$$

- **Sigmoid** The Sigmoid activation takes in the feature x and applies the following function element-wise:

$$y = \sigma(x) = \frac{1}{1 + \exp(-x)} \quad (7)$$

The gradient of Sigmoid is:

$$\frac{\partial y}{\partial x} = \sigma(x)(1 - \sigma(x)) \quad (8)$$

- **SELU** The SELU activation takes in the feature x and applies the following function element-wise:

$$y = \text{SELU}(x) = \text{scale} * \begin{cases} x & \text{if } x > 0 \\ \alpha e^x - \alpha & \text{if } x \leq 0 \end{cases} \quad (9)$$

where $\alpha = 1.6732632423543772848170429916717$, $\text{scale} = 1.0507009873554804934193349852946$. The gradient of SELU is:

$$\frac{\partial y}{\partial x} = \text{scale} * \begin{cases} 1 & \text{if } x > 0 \\ \alpha e^x & \text{if } x \leq 0 \end{cases} \quad (10)$$

The final returned gradient of these activation functions is $\frac{\partial \ell}{\partial x} = \frac{\partial \ell}{\partial y} \odot \frac{\partial y}{\partial x}$, where \odot is the element-wise product, and will be propagated to the former layer.

4) *Sequential*: A *Sequential* container takes in a list of layers as input and combines them together in order. In forward function, each layer's own forward function is called sequentially to form a complete forward propagation. In backward function, each layer's own backward function is called in reverse order sequentially to form a complete backward propagation.

5) *MSELoss*: *MSELoss* module computes the **Mean Squared Error Loss** between the model output y and the target t :

$$\text{loss} = ||y - t||^2 \quad (11)$$

The final loss value is the mean value of the batch size samples. The returned gradient is $\frac{\partial \ell}{\partial y} = 2(y - t)$, averaged by the number of elements in the target tensor.

6) *SGD Optimizer*: The SGD optimizer updates the model parameters with the following rule:

$$x_{t+1} = x_t - \gamma \frac{\partial \ell}{\partial x_t} \quad (12)$$

where γ is the learning rate.

In our implementation, SGD optimizer provides two interfaces for the end-user: `zero_grad()` and `step()`, which performs the same functions as the implementation in PyTorch. The `zero_grad()` clears all the gradients of the parameters

in the model and the `step()` performs the SGD update step for all the parameters.

For loss backward propagation, there are some differences between our framework and PyTorch. In our framework, the computed loss can not perform backward propagation. Instead, the `criterion` object, which is an instance of the `MSELoss`, performs the backward propagation. Besides, we need to perform the model backward propagation manually. These differences are shown below.

```
1 optimizer.zero_grad() # zero the gradient buffers
2 output = model(input)
3 loss = criterion(output, target)
4 loss.backward() # loss backward
5 optimizer.step() # parameter update
```

Listing 2: PyTorch Loss Backward Propagation

```
1 optimizer.zero_grad() # zero the gradient buffers
2 output = model(input)
3 loss = criterion(output, target)
4 dldy = criterion.backward() # loss backward
5 model.backward(dldy) # model backward
6 optimizer.step() # parameter update
```

Listing 3: Our Framework Loss Backward Propagation

7) *MLP model*: The MLP model used for binary classification consists of 3 hidden layers, each with 25 units. The input dimension is 2 while the output dimension is 1. Using the above implemented modules, the model definition in our framework is the same as in PyTorch, using a *Sequential* container to wrap the *Linear* layers and activation layers.

II. EXPERIMENT RESULTS

To test our framework, we generate a training and test set of 1000 points sampled uniformly in $[0, 1]^2$. To each of them, we associate the label 0 if the point is outside the disk of radius $\frac{1}{\sqrt{2\pi}}$, and 1 otherwise. Therefore, the objective of the model is to predict if a given point is outside or inside the disk, a binary classification problem.

A. Effectiveness of Adding Sigmoid

We implement 4 activation functions: ReLU, Tanh, SELU and Sigmoid. We find that a MLP model with all Sigmoid activation functions cannot be successfully trained, even with the standard PyTorch functions. For binary classification tasks, usually a Sigmoid activation functions will be applied to the output. We verified the necessity of adding Sigmoid at the end by experiments. Each models are trained for 300 epochs, with batch size 100 and learning rate 0.1. The training losses are shown in Fig. 2.

From this figure, we can see that adding Sigmoid activation function at the end of the model smooths the training loss curves and results in lower training losses.

B. Comparison of Different Activation Functions

From Fig. 2 we can also see that models with ReLU and SELU have lower training losses than the model with Tanh. This can also be verified by the final test accuracy, as shown in Fig. 3. We train and test the models with ReLU, Tanh,




Fig. 2: The training loss of different activation (combinations).

SELU for 10 rounds, with batch size 100 and learning rate 0.1. These models achieve the average test accuracy 98.64%, 97.36%, 97.54% and the test accuracy standard deviation 0.0049, 0.0061, 0.0116 respectively.




Fig. 3: The test accuracy of different activation.

III. SUMMARY

In this project, we implement a mini deep learning framework, which supports Sequential models, and provides Linear, ReLU, Tanh, SELU, Sigmoid activations, MSELoss function and SGD optimizer.

We build a simple MLP model using this framework, to verify the effectiveness of this framework. Experiments show that our framework can be used to implement and train the model successfully and can achieve a high performance on a simple binary classification task.

However, there are still some limitations of our framework. We only implement MSELoss, while for binary classification tasks, the BCELoss is preferred.