

Deep Learning Project1 Report

Chenkai Wang
323452

chenkai.wang@epfl.ch

Siran Li
321825

siran.li@epfl.ch

Shijian Xu
327448

shijian.xu@epfl.ch

Abstract—The objective of this project is to implement a mini “deep learning framework” using only PyTorch’s tensor operations and the standard math library. The task is to use this mini framework to train a Multi-Layer Perceptron (MLP) for a binary classification problem. To achieve this goal, we implement the modules Linear, ReLU, Tanh, Sequential, MSELoss and SGD optimizer. The interface of these modules are very close to their PyTorch implementation.

I. IMPLEMENTATION

The mini deep learning framework consists of several modules: Linear, Sequential, MSELoss, SGD optimizer and several activation functions, which are ReLU, Tanh, Sigmoid and SELU. All of these modules, except the SGD optimizer, inherit from a basic Module class. The final MLP model also inherits from the Module class and is composed of a Sequential member.

A. Modules

1) *Module*: The base class Module defines the functions each module needs to implement. The code snippet of the implementation of the basic Module class is shown below.

```
1 class Module(object):
2     """
3     Base Class. Other modules should inherit from it
4     and rewrite the 'forward', 'backward' and
5     'param' functions.
6     """
7     def __call__(self, *input):
8         return self.forward(*input)
9
10    def forward(self, *input):
11        raise NotImplementedError
12
13    def backward(self, *gradwrtoutput):
14        raise NotImplementedError
15
16    def parameters(self):
17        return []
```

Listing 1: The Base Module Class

We also use the Python built-in function `__call__` to make the modules’ interfaces behave more PyTorch-like. In this case, we can call the module directly (`model(x)`) instead of `model.forward(x)`.

2) *Linear*: Linear layer is the core in the framework. In our implementation, this layer takes in a 2-d tensor $x \in \mathbb{R}^{N \times in_features}$, where N is the batch size. The forward output is produced by $y = xW + b$, where W is the weight matrix with shape $in_features \times out_features$, b is the bias vector. Both W and b are initialized with uniform distribution from $\mathcal{U}(-\sqrt{k}, \sqrt{k})$, where $k = \frac{1}{in_features}$. The backward

function takes in the gradient of loss with respect to the output y , $\frac{\partial l}{\partial y}$, and compute the gradient for W, b and x : $\frac{\partial l}{\partial W} = x^T \frac{\partial l}{\partial y}$, $\frac{\partial l}{\partial b} = \frac{\partial l}{\partial y}$, $\frac{\partial l}{\partial x} = \frac{\partial l}{\partial y} W^T$. The gradient of x will be returned and propagated to the former layer.

3) Activation Functions:

- ReLU The ReLU activation takes in the feature x and applies the following function element-wise:

$$y = ReLU(x) = (x)^+ = \begin{cases} x & \text{if } x > 0 \\ 0 & \text{if } x \leq 0 \end{cases} \quad (1)$$

The gradient of ReLU is:

$$\frac{\partial y}{\partial x} = \begin{cases} 1 & \text{if } x > 0 \\ 0 & \text{if } x \leq 0 \end{cases} \quad (2)$$

- Tanh The Tanh activation takes in the feature x and applies the following function element-wise:

$$y = Tanh(x) = \frac{\exp(x) - \exp(-x)}{\exp(x) + \exp(-x)} \quad (3)$$

The gradient of Tanh is :

$$\frac{\partial y}{\partial x} = 1 - Tanh(x)^2 \quad (4)$$

- Sigmoid The Sigmoid activation takes in the feature x and applies the following function element-wise:

$$y = \sigma(x) = \frac{1}{1 + \exp(-x)} \quad (5)$$

The gradient of Sigmoid is:

$$\frac{\partial y}{\partial x} = \sigma(x)(1 - \sigma(x)) \quad (6)$$

- SELU The SELU activation takes in the feature x and applies the following function element-wise:

$$y = SELU(x) = scale * \begin{cases} x & \text{if } x > 0 \\ \alpha e^x - \alpha & \text{if } x \leq 0 \end{cases} \quad (7)$$

where $\alpha = 1.6732632423543772848170429916717$, $scale = 1.0507009873554804934193349852946$. The gradient of SELU is:

$$\frac{\partial y}{\partial x} = scale * \begin{cases} 1 & \text{if } x > 0 \\ \alpha e^x & \text{if } x \leq 0 \end{cases} \quad (8)$$

The final returned gradient is $\frac{\partial l}{\partial x} = \frac{\partial l}{\partial y} \frac{\partial y}{\partial x}$ and will be propagated to the former layer.

4) *Sequential*: A *Sequential* container takes in a list of layers as input and combines them together in order. In forward function, the layers' own forward functions are called sequentially to form a complete forward propagation. In backward function, the layers' own backward functions are called in reverse order sequentially to form a complete backward propagation.

5) *MSELoss*: *MSELoss* module computes the Mean Squared Error between the model output y and the target t :

$$\text{loss} = ||y - t||^2 \quad (9)$$

The final loss value is the mean value of the batch size samples. The returned gradient is $\frac{\partial l}{\partial y} = 2(y - t)$, averaged by the number of elements in the target tensor.

6) *SGD Optimizer*: The *SGD* optimizer updates the model parameters with the following rule:

$$x_{t+1} = x_t - \gamma \frac{\partial l}{\partial x_t} \quad (10)$$

where γ is the learning rate.

In our implementation, *SGD* optimizer provides two interfaces: *zero_grad* and *step*, which performs the same functions as the implementation in *PyTorch*. For loss backward propagation, there are some differences between our framework and *PyTorch*.

```
1 optimizer.zero_grad() # zero the gradient buffers
2 output = model(input)
3 loss = criterion(output, target)
4 loss.backward() # loss backward
5 optimizer.step() # parameter update
```

Listing 2: *PyTorch* Loss Backward Propagation

```
1 optimizer.zero_grad() # zero the gradient buffers
2 output = model(input)
3 loss = criterion(output, target)
4 dldy = criterion.backward() # loss backward
5 model.backward(dldy) # model backward
6 optimizer.step() # parameter update
```

Listing 3: Our Framework Loss Backward Propagation

7) *MLP model*: The *MLP* model used for binary classification consists of 3 hidden layers, each with 25 units. The input dimension is 2 while the output dimension is 1. Using the above implemented modules, the model definition in our framework is the same as in *PyTorch*, using a *Sequential* container to wrap the *Linear* and activation layers.

II. EXPERIMENT RESULTS

To test our framework, we generate a training and test set of 1000 points sampled uniformly in $[0, 1]^2$. To each of them, we associate the label 0 if the point is outside the disk of radius $\frac{1}{\sqrt{2\pi}}$, and 1 otherwise. Therefore, the objective of the model is to predict if a given point is outside or inside the disk, a binary classification problem.

A. Effectiveness of Adding Sigmoid

We implement 4 activation functions: *ReLU*, *Tanh*, *SELU* and *Sigmoid*. It is strange to see that the model with all sigmoid activation can not be trained. Why ??? For binary classification tasks, usually a *Sigmoid* activation functions will be applied to the output. We verified the necessity of adding *Sigmoid* at the end by experiments. Each models are trained for 300 epochs, with batch size 100 and learning rate 0.1. The training losses are shown in Fig. 1.




Fig. 1: The training loss of different activation (combinations).

From this figure, we can see that adding *Sigmoid* activation function at the end of the model smooths the training loss curves and results in lower training losses.

B. Comparison of Different Activation Functions

From Fig. 1 we can also see that models with *ReLU* and *SELU* have lower training losses than the model with *Tanh*. This can also be verified by the final test accuracy, as shown in Fig. 2. We train and test the models with *ReLU*, *Tanh*, *SELU* for 10 rounds, with batch size 100 and learning rate 0.1.




Fig. 2: The training loss of different activation (combinations).

III. SUMMARY

There are some limitations in our framework. For binary classification tasks, *MSELoss* is not a good choice. It is better to use the binary cross entropy loss.